

NATURAL LANGUAGE PROCESSING

WITH DEEP LEARNING



University of Guilan

Lecturer: Javad PourMostafa

NLP981

Introduction to NLP

MAIN APPROACHES

1. Rule-based Methods
 1. Regular Expression
 2. CFG
 3. ...
2. Traditional ML algorithms and Probabilistic modeling
 1. Likelihood Maximization
 2. Supervised/Non-supervised algorithms
 3. ...
3. Deep Learning
 1. RNN
 2. CNN
 3.

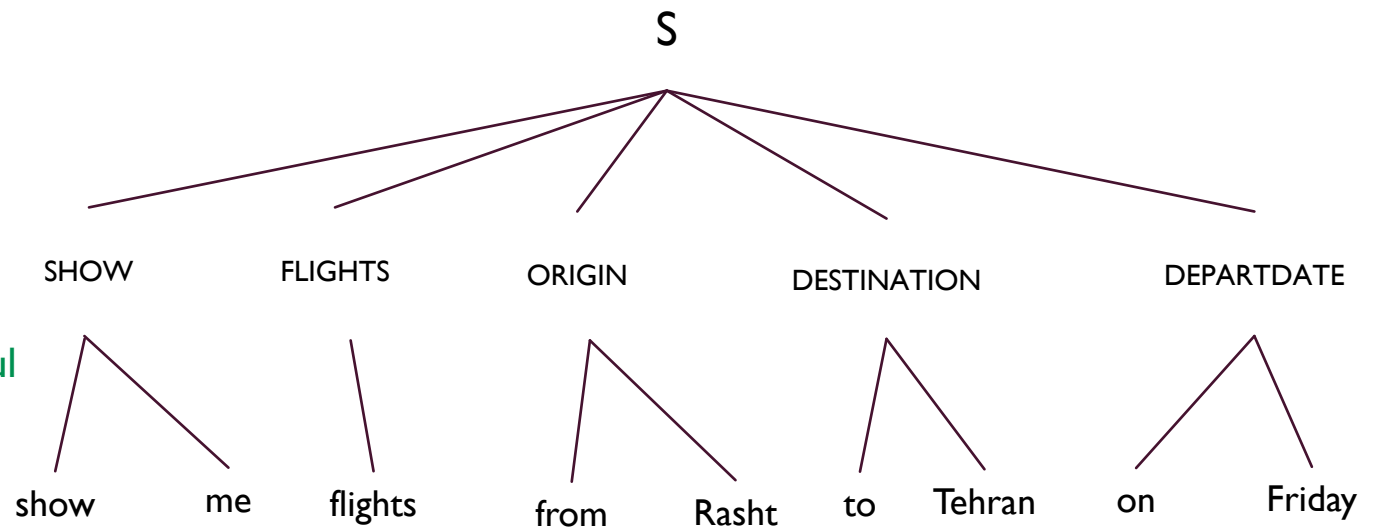
SEMANTIC SLOT FILLING: CFG

■ What are Context Free Grammar?

- A formal grammar
- Every production rule is of the form $V \rightarrow w$
- V is a single nonterminal symbol and w is a string of terminals and/ or non-terminals

■ Context Free Grammar:

- $SHOW \rightarrow \text{show me} \mid \text{i want} \mid \text{can i see} \mid \dots$
- $FLIGHTS \rightarrow (\text{a}) \text{ flight} \mid \text{flights}$
- $ORIGIN \rightarrow \text{from CITY}$
- $DESTINATION \rightarrow \text{to CITY}$
- $CITY \rightarrow \text{Rasht} \mid \text{Tehran} \mid \text{Dubai} \mid \text{Isfahan} \mid \text{Istanbul}$



SEMANTIC SLOT FILLING: CFG (CONT.'S)

- Advantages vs. Disadvantages?
 - Usually done manually
 - Time Consuming
 - High Precision
 - Low Recall

CONFUSION MATRIX

- Machine Learning Fundamental
- Dividing data into Training set and Testing set
- Summarize performance on the Testing data
- Create confusion matrix

		Actual Values	
		Positive (1)	Negative (0)
Predicted Values	Positive (1)	TP	FP
	Negative (0)	FN	TN

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$\text{Precision} = \frac{\text{True Positive}}{\text{True Positive} + \text{False Positive}}$$

$$\text{Recall} = \frac{\text{True Positive}}{\text{True Positive} + \text{False Negative}}$$

SEMANTIC SLOT FILLING: CRF

- A Probabilistic Graphical Model
- Training Corpus:

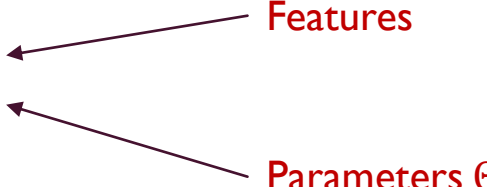
ORIG DEST DATE

Show me flights from Rasht to Tehran on friday.

- Some Feature Engineering:
 - Is the word in a list of city names?
 - What is the previous slot?
 - Is the word capitalized?
 -

SEMANTIC SLOT FILLING: CRF (CONT'D)

- Define your Conditional Random Field model:

$$P(\text{tags} \mid \text{words}) = \dots$$


Features

Parameters Θ

- Training:


$$P(\text{tags} \mid \text{words}) \longrightarrow \text{Max by } \Theta$$

- Inference/Evaluation/Test:

$$\text{tags}^* = \operatorname{argmax} P(\text{tags} \mid \text{words})$$

SEMANTIC SLOT FILLING: LSTM


- Deep Learning Approach
- Long-term Sequences
- No feature generation
- Defining Model, Train and Inference



DEEP LEARNING VS. TRADITIONAL NLP

- WHY DL?
 - State-of-the-art performance in subtasks
 - Sentiment Analysis
 - MT
 - Look fancy and has upward trend
 - Most research is happening (ACL, EMNLP, arxiv, etc.)
- Our strategy?
 - Mainly work on DL approaches
 - However study traditional ones

NLP PYRAMID



MORPHOLOGY

- Analyses how words formed
- What is their origin
- Mostly deal with:
 - Prefixes/suffixes
 - Gender detection
 - Lemmatization
 - Spell checking
- Operations are at word level (viewed as a sequence of chars)

SYNTAX

- Underlying structure of a sentence
- Refer to grammar
- Most researched branch of CL
- Tasks:
 - POS tagging
 - Building Syntax Trees
 - Building Dependency Trees
- Usually works on sentences (viewed as sequence of words)

SEMANTIC

- Derives meaning from text
- Known problems
 - Name Entity Recognition
 - Relation Extraction
 - Semantic Role Labeling (shallow semantic parsing)
 - Word Sense Disambiguation
- Works on sentences (viewed as a sequence of words)

PRAGMATICS

- Analyses the text as a whole
- Tasks:
 - Summarization
 - Topic Segmentation
 - Coreference / Anaphora resolution (find out what word refers what.)
- Works on a text (viewed as a sequence of sentences)

Text Classification

SCENARIO OF TEXT CLASSIFICATION

- Task : Sentiment Analysis
- Input: IMDB movie reviews dataset
- Output: Polarity
- Pos example:
 - “ I've seen this story before but my kids haven't.
Boy with troubled past joins military, faces his past, falls in love and becomes a man. ”
- Neg example:
 - “ I feel about DARLING LILI.This massive musical is so peculiar and over blown, over produced and must have ”

WHAT'S TEXT?

- A sequence of
 - Characters
 - ✓ **Words**
 - Phrases
 - Sentences
 - Paragraphs
 - ...
- Seems natural to think of a text as a sequence of words!
- How to find the boundaries of words?
 - Punctuation or spaces

TOKENIZATION

- Task of chopping given a defined doc unit up into pieces, called **tokens**.
- A good example of simple whitespace tokenizer in python:
 - NLTK package for English :
 - `nltk.tokenize.WhitespaceTokenizer`,
 - `nltk.tokenize.TreebankWordTokenizer()`
 - `nltk.tokenize.WordPuncTokenizer()`
 - Hazm package for Persian:
 - `word_tokenize()`

TOKEN NORMALIZATION

- Stemming: Find the root form of the word called the **stem**
 - by removing and replacing suffixes
 - Porter's stemmer
- Lemmatization: Return the base of dictionary form of a word known **lemma**
 - WordNet lemmatizer
- Like:
 - wolves → wolf
 - talks → talk

PORTER'S STEMMER

- An algorithm for suffix stripping
- Using 5 heuristic phases for word reduction
- Applied sequentially
- Example of phase one rules:
 - SSES → SS caresses → caress
 - IES → I ponies → poni
 - SS → SS. caress → caress
 - S → cats → cats
- Problem with irregular forms, produce non-words

nlTK.stem.PorterStemmer examples:

-feet → feet
-wolves → wolv
-cats → cats
-talked → talk

WORDNET LEMMATIZER

- Uses the WordNet database to lookup lemmas
- Only removes affixes if the resulting word is in its dic
- `nltk.stem.WordNetLemmatizer`
- **Not all forms are reduced!**

`nltk.stem.WordNetLemmatizer` examples:

-feet → foot
-wolves → wolf
-cats → cat
-talked → talked

OTHER NORMALIZATION

- Normalization capital letters
 - approaches:
 - Lowercasing the beginning of the sentence
 - Lowercasing words in titles
 - Leave mid-sentence words as they are
 - Use ML to retrieve **true casing**
- Acronyms
 - Frequently use in text msg and chats
 - eta, e.t.a., E.T.A. → E.T.A.
 - Write regular expressions, however it's hard!

BAG OF WORDS (BOW)

- Count the occurrence of a specific token in our text
- For each token we'll have a feature column called

- **Text Vectorization**

- Example:

- Sentence 1: I like to play tennis
 - Sentence 2: Did you go outside to play tennis
 - Sentence 3: John and I play tennis

Sentences	like	play	tennis	go	outside
#1	1	1	1	0	0
#2	0	1	1	1	1
#3	0	1	1	0	0

- Disadvantages?

- Careful design in vocabulary
 - Sparsity
 - Discard word order in the context (This is interesting vs. Is this interesting?)

PRESERVE SOME ORDERING: N-GRAMS

- Solution: using **n-grams** technique
- Like: 2-grams for token pairs and so forth.

- Example:

- Sentence 1: I like to play tennis
- Sentence 2: Did you go outside to play tennis
- Sentence 3: John and I play tennis

Sentences	like to	play tennis	tennis	go outside	outside
#1	1	1	1	0	0
#2	0	1	1	1	1
#3	0	1	1	0	0

- Disadvantages?

- Too many text vectorizations

REMOVE SOME N-GRAMS

- There are 3 types of n-grams based on their occurrence in documents:
- **High Frequency**
 - Articles, prepositions, etc.
 - Called **Stop-words** → Remove them
- **Low Frequency**
 - Typos and rare n-grams → remove them
 - If we like to **overfit** our model, then it will be a good choice
- **Medium Frequency**
 - Good n-grams → Sustainable

TF-IDF

- **There are numerous medium frequency n-grams**
- Filtering out bad n-grams is useful
- Would be better if we set a **ranking system for medium frequency n-grams**
- **Idea:** n-grams with smaller frequency can be more discriminating
- **Cause:** capture a specific issue in the review

TERM FREQUENCY (TF)

- $tf(t, d)$
- Frequency for term or n-gram t in the document d
- Variants:

Weighting scheme	TF weight
binary	0, 1
Raw count	$f_{t,d}$
Term frequency	$f_{t,d} / \sum_{t' \in d} f_{t',d}$
Log normalization	$1 + \log(f_{t,d})$

INVERSE DOCUMENT FREQUENCY (IDF)

- $N = |D|$ - total num of documents in corpus
- $|\{d \in D : t \in d\}|$ - num of documents where term t appears
- $idf(t, D) = \log \frac{N}{|\{d \in D : t \in d\}|}$
- **TF-IDF:**
 - $tfidf(t, d, D) = tf(t, d).idf(t, D)$
- **Achievement:**

We will find high frequent term in the given document
but low frequent term in the whole collection of documents.

TF-IDF EXAMPLE

- Doc 1: I like to play tennis
- Doc 2: Did you go outside to play tennis
- Doc 3: John and I play tennis outside

- **Replace counters with TF-IDF values**

- $\text{tf-idf}(\text{like to}, \text{doc 1})$:

- $1/3 \times \log(3/1) =$
- $0.33 \times 0.47 \cong 0.15$


- $\text{tf-idf}(\text{outside}, \text{doc 2})$:

- $1/3 \times \log(3/2) =$
- $0.33 \times 0.176 \cong 0.058$

documents	like to	play tennis	tennis	outside
#1	0.15	--	--	--
#2	--	--	--	0.05
#3	--	--	--	--

LINEAR MODEL FOR CLASSIFICATION

- Task: Sentiment Classification
- Dataset: IMDB movie reviews dataset
- http://ai.stanford.edu/~amaas/data/sentiment/acllmbd_v1.tar.gz



- 25000 reviews from IMDB
- Contain an **even** num of pos and negative reviews
- Randomly guessing yields, how many accuracy?
 - 50 percent
- A **negative** review has a score ≤ 4 : label **0** and A **positive** review has a score ≥ 7 : label **1**

BINARY LOGISTIC REGRESSION


- Data split: 50/50
- Evaluation metric: accuracy
- Features: bag of 1-grams with TF-IDF values
- Extremely sparse feature matrix (25K rows and 74849 columns for TS)
- 99.8% are zeros
- Suggested Model for training: **Logistic Regression**
 - In statistics, uses to model probabilities of a certain class
 - Lose/win, pass/fail, alive/dead
 - **Linear** classification model
 - Handle sparse data
 - Weights can be interpreted

SIGMOID ACTIVATION FUNCTION

- In order to **map** predicted values to probabilities
- Maps any real values into another value between 0 and 1


$$■ S(\zeta) = \frac{1}{1 + e^{-\zeta}}$$

- $S(\zeta)$ = output between 0 and 1 (probability estimate)
 - ζ = input to the function (your algorithm's prediction e.g. $mx + b$)
 - e = base of natural log
-
- In order to map discrete class, we select a **threshold** like:
 - $P \geq 0.5, class = 1$ and $P < 0.5, class = 0$



GRADIENT DESCENT

- An optimization algorithm
- Find the parameters that minimize cost functions
- Parameters mean **coefficients** in linear regression and **weights** in neural networks
- In ML use it to **update the parameters of our model.**
- Iteratively moving in the direction of **steepest descent**
- Learning rate = step size
- E.g.: Given cost function:
 - $$f(m, b) = \frac{1}{2n} \sum_{i=1}^n (y_i - (mx_i + b))^2$$
- Gradient can be calculated from the **partial derivatives.**



GRADIENT DESCENT EXAMPLE

- We need to update **m** and **b** called **weights**
- $y = mx + b$
- $h(x) = mx + b$
- $h(x) = y$, $h(x)$ called hypothesis in ML but this y is not actual value
- This is predicted y from our hypothesis
- Assumption: $b = 1$ and $m = 0.5$ and $x = 10$
- Then $h(x) = 0.5 \times 10 + 1$
- *predicted y is 6 but actual value is 5*
- $error = (h(x) - y)^2 = (6 - 5)^2 = 1$
- Expo 2 used to get rid of negative values
- Repeat for all data points in our DS and sum up all of them called **cost function**

X	Y
10	5
12	6.6
12	6.6
3	1
---	---

BETTER IMPROVEMENTS

- Add **2-grams** feature extraction (156821 columns)
- Throw away n-grams seen less than 5 times (min_frequency)
- Use special token like emoticons (;-))
- Adding stemming and lemmatization
- Apply different models
- Instead of using TF-IDF, use deep learning for word embedding

ANOTHER SCENARIO

- N-gram \rightarrow feature index
- What if, when there are **2TB** of texts?
- Problem on distributed systems
- Hold a very large vocabulary dictionary in memory
- So what's the solution?
 - Using **hashing** trick for squeezing
 - **N-gram \rightarrow hash(n-gram) % 2^b**
 - Related paper: Feature Hashing for Large Scale Multitask Learning
 - <https://arxiv.org/pdf/0902.2206.pdf>

HASHING FEATURES

- Using a hashing function
- Maps any n-grams to a number range
- No need to save n-grams in a dictionary
- Might map any n-grams to the same number
 - Collision
- The larger range then less chance of collisions

