

NATURAL LANGUAGE PROCESSING

WITH DEEP LEARNING



University of Guilan

Lecturer: Javad PourMostafa

NLP981

BOW VECTORIZATION

- **BOW** is a **sum of sparse one-hot-encoded vectors** corresponding to each particular word.

	I	Like	Play	Tennis	Outside	Enjoy
I	1	0	0	0	0	0
	+					
like	0	1	0	0	0	0
	+					
tennis	0	0	0	1	0	0
	=					
I like tennis	1	1	0	1	0	0

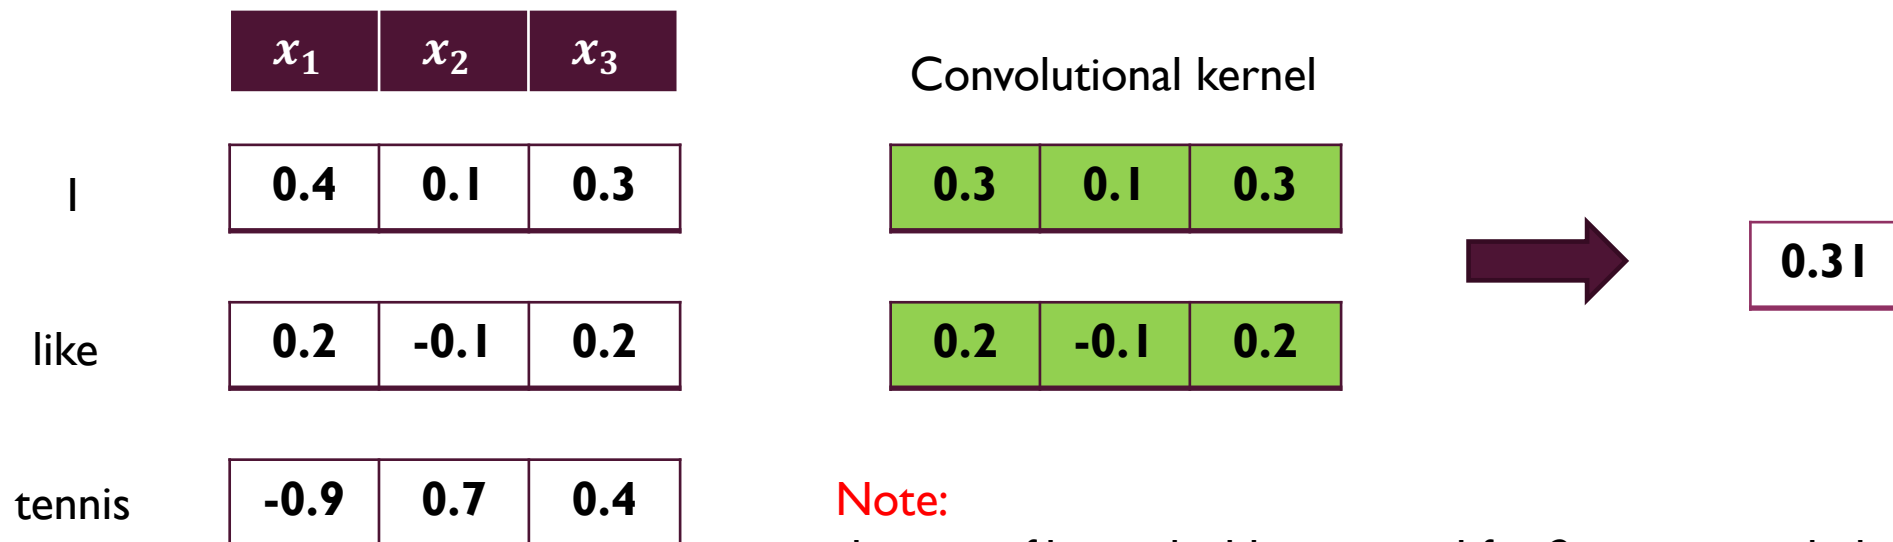
NEURAL VECTORIZATION

- Dense-based method
- Less columns (about 300)
- Pre-trained and online NN
- A good example of pre-trained is **word2vec embedding**
- Word2vec property:
 - **Words that have similar context tend to have collinear vectors**

	x_1	x_2	x_3
I	0.4	0.1	0.3
	+		
like	0.2	-0.1	0.2
	+		
tennis	-0.9	0.7	0.4
	=		
	-0.3	0.7	0.9

1-D CONVOLUTIONAL LAYER

- How do we make 2-grams?
- **Don't have Word2vec embedding for token pairs**
- **Use convolutional layer**



Note:

this conv filter is highly activated for 2-grams entitled
“preferences in racket-based sports”

LANGUAGE MODELING

- This is the

- **The house**


- **Rat**

- **Did**

- **malt**

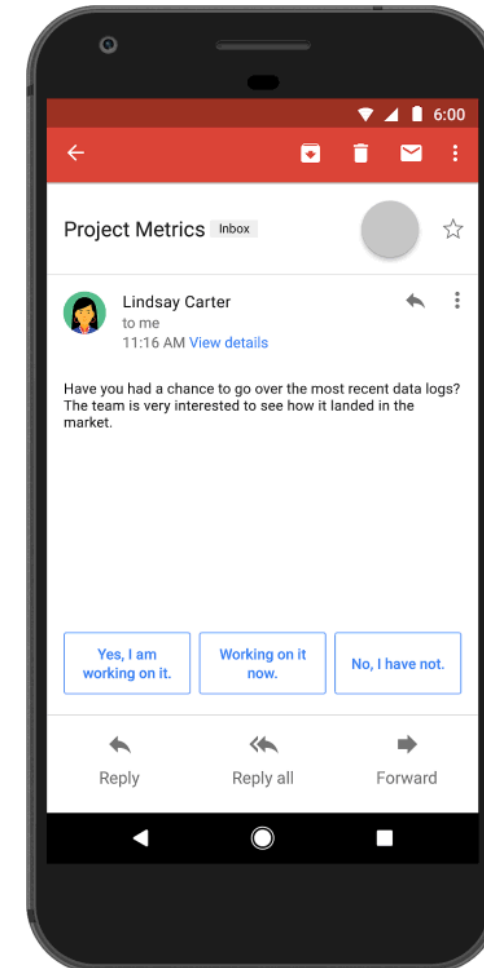
- The probability of the next word?

- $P(\text{the house} \mid \text{this is the}) = ?$



WHERE DO WE NEED LANGUAGE MODELING?

- Smart Reply
- Suggestion in messengers
- Spelling correction
- Machine translation
- Speech recognition
- Handwriting recognition
- ...



TOY CORPUS, 4-GRAMS

- This is the house that Jack built
- This is the malt
- That lay in the house that jack built
- This is the rat
- That ate the malt
- That lay in the house that jack built
- This is the cat
- That killed the rat
- That ate the malt
- That lay in the house that jack built

- $P(\text{house} \mid \text{this is the}) =$

- $\frac{C(\text{this is the house})}{C(\text{this is the ...})} =$

- $\frac{1}{4}$

TOY CORPUS, BI-GRAMS

- This is the house that Jack built
- This is the malt
- That lay in the house that jack built
- This is the rat
- That ate the malt
- That lay in the house that jack built
- This is the cat
- That killed the rat
- That ate the malt
- That lay in the house that jack built

- $P(Jack \mid that) =$

- $\frac{C(that \ Jack)}{C(that \dots)} =$

- $\frac{4}{10}$

PROBABILITY OF A SEQUENCE OF WORDS

- $W = (w_1, w_2, w_3, \dots, w_n)$, $P(w_5 \mid w_1, w_2, w_3, w_4)$

- Chain Rule:

$$P(w_1^n) = P(w_1) P(w_2 \mid w_1) P(w_3 \mid w_1^2) \dots P(w_n \mid w_1^{n-1}) = \prod_{k=1}^n P(w_k \mid w_1^{k-1})$$

- Example:

$$P(\text{"Its water is so transparent"}) =$$

$$P(\text{its}) \times$$

$$P(\text{water} \mid \text{its}) \times$$

$$P(\text{is} \mid \text{its water}) \times$$

$$P(\text{so} \mid \text{its water is}) \times$$

$$P(\text{transparent} \mid \text{its water is so})$$

PROBABILITY OF A SEQUENCE OF WORDS (CONT'D)

- Last term is too long, how to deal with that?!
- Solution?
 - Markov Assumption and intuition of the n-gram model
- “According to Markov property , we can predict the probability of some future state without looking too far in the past”.

$$p(w_i | w_1 \dots w_{i-1}) = p(w_i | w_{i-n+1} \dots w_{i-1})$$

- Should not care about all the history!
- So Just approximate the history by the last few words

BIGRAM MODEL

- Approximates by using **only** the conditional probability of the **preceding word**
- **E.g.**
- $P(\text{the} \mid \text{Walden Pond's water is so transparent that}) \approx P(\text{the} \mid \text{that})$

$$P(W_n \mid W_1^{n-1}) \approx P(W_n \mid W_{n-1})$$

- Main Idea: Markov assumption

GENERALIZE THE BIGRAM TO N-GRAM

- We can generalize bigram to the **trigram**
 - **Which looks two words into the past**
- Thus to the **n-gram**
 - **Which looks n-1 words into the past**
 - Thus the **general** equation

$$P(W_n | W_1^{n-1}) \approx P(W_n | W_{n-N+1}^{n-1})$$

- **We can compute the probability of a complete word sequence (bigram model)**

$$P(W_1^n) \approx \prod_{k=1}^n P(W_k | W_{k-1})$$

MAXIMUM LIKELIHOOD ESTIMATION (MLE)

- How do we estimate these bigram or n-gram probabilities?
 - By **MLE**
- 1. **Get the MLE estimate** for the parameters of an n-gram model
 - How? **Counts from corpus**
- 2. **Normalizing** the counts
 - Lie between **0 and 1**, How?
 - **Dividing by some total count**

$$P(w_n|w_{n-1}) = \frac{C(w_{n-1}w_n)}{C(w_{n-1})} \text{ For bigram model}$$

EXAMPLE OF BIGRAM USING MINI-TRAINING CORPUS

- $\langle S \rangle$ I am Sam $\langle /S \rangle$
- $\langle S \rangle$ Sam I am $\langle /S \rangle$
- $\langle S \rangle$ I do not like green eggs and ham $\langle /S \rangle$
- Add fake tokens ($\langle S \rangle$ and $\langle /S \rangle$)
 - $\langle S \rangle$ gives us the bigram context of the first word.
 - $\langle /S \rangle$ we need a true probability distribution.
- Some bigram probabilities (**relative frequency**):

$$\begin{array}{lll} P(I | \langle s \rangle) = \frac{2}{3} = .67 & P(\text{Sam} | \langle s \rangle) = \frac{1}{3} = .33 & P(\text{am} | I) = \frac{2}{3} = .67 \\ P(\langle /s \rangle | \text{Sam}) = \frac{1}{2} = 0.5 & P(\text{Sam} | \text{am}) = \frac{1}{2} = .5 & P(\text{do} | I) = \frac{1}{3} = .33 \end{array}$$

GENERATIVE MODEL

- An N-gram model can be seen as a probabilistic **automata** for generating sentences.
- How?
 - Initialize sentence with N-1 $\langle s \rangle$ symbols
 - Until $\langle /s \rangle$ is generated do:

Stochastically pick the next word based on the conditional probability of each word given the previous N - 1 words.

$$\hat{\lambda} = \operatorname{argmax}_{\lambda} P(T \mid M(\lambda))$$

EVALUATING LANGUAGE MODELS

- Ideally, evaluate use of model in end application (**extrinsic, in vivo**)
 - Realistic
 - Often very expensive
- Evaluate on ability to model test corpus (**intrinsic**)
 - Less realistic
 - Cheaper

INTRINSIC EVALUATION

Training Data

i live in osaka
i am a graduate student
my school is in nara
...

Train
Model

Model

Testing Data

i live in nara
i am a student
i have lots of homework
...

Test
Model

Model Accuracy

Perplexity

PERPLEXITY (PP)

- Call **hold-out (data) perplexity**
- Is the inverse probability of the test set
- **Measure of how well a model “fits” the test data**
- **Lower perplexity is better** because the greater conditional probability is better
- Perplexity of W (test set) with a bigram model: (N is the number of words in test sentence)

$$PP(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i|w_{i-1})}}$$

PERPLEXITY (PP) EXAMPLE

- Train Corpus:
 - This is the house that Jack built.
- Test Corpus:
 - This is the **malt**.
- Perplexity of bigram language model?

$$P(\text{malt} \mid \text{the}) = \frac{C(\text{the malt})}{C(\text{the})} = 0$$

$$P(W \text{ test data}) = 0$$

PP = infinite

UNKNOWN OR OOV

- The percentage of OOV words that appear in test set
 - **OOV rate**
- Solution?
 - Build a vocabulary
 - How? By min-frequency
 - Substitute OOV words with pseudo-word called <UNK> (train and test)
 - Compute counts as usual for all tokens

PERPLEXITY WITHOUT OOV WORDS

- Train Corpus:
 - This is the house that Jack built.
- Test Corpus:
 - This is Jack.
- Perplexity of bigram language model?

$$P(Jack | is) = \frac{C(is Jack)}{C(is)} = 0$$

$$P(W \text{ test data}) = 0$$

PP = infinite

SMOOTHING OR DISCOUNTING

- What do we do with words that are in our vocabulary but appear in a test set in an unseen context?
- Stop zero probability by a modification
- Ways to do smoothing:
 1. **Add-one smoothing**
 2. **Add-k smoothing**
 3. **Stupid backoff**
 4. **Kneser-Ney smoothing**

LAPLACE SMOOTHING

- Idea:
 - **Add one** to all the **bigram counts**, before we normalize them into probabilities.
- Then, all the counts that used to be zero will have a count of 1, the counts of 1 will be 2 and so on.
- Laplace does not work well in n-gram models.
- But a good baseline for other smoothing methods

LAPLACE SMOOTHING (CONT'D)

- Recall **unsmoothed maximum likelihood estimate of unigram** probabilities:

$$P(w_i) = \frac{c_i}{N}$$

- **Add-one smoothing:**

- Just adds one to each count

$$P_{Laplace}(w_i) = \frac{c_i + 1}{N + V}$$

- Add-one smoothing **bigram** probability:

$$P_{Laplace}(W_n | W_{n-1}) = \frac{C(w_{n-1}w_n) + 1}{C(w_{n-1}) + V}$$

ADD-ONE SMOOTHING EXAMPLE

	i	want	to	eat	chinese	food	lunch	spend
i	6	828	1	10	1	1	1	3
want	3	1	609	2	7	7	6	2
to	3	1	5	687	3	1	7	212
eat	1	1	3	1	17	3	43	1
chinese	2	1	1	1	1	83	2	1
food	16	1	16	1	2	5	1	1
lunch	3	1	1	1	1	2	1	1
spend	2	1	2	1	1	1	1	1

Add-one smoothed bigram counts for eight of the words (out of $V = 1446$) in the Berkeley Restaurant Project corpus of 9332 sentences. Previously-zero counts are in gray.

ADD-K SMOOTHING

- Idea:
 - Pull some probabilities from frequent bigram and to infrequent ones
- Tune k with **test data**

$$P_{Add-k}(w_n | w_{n-1}) = \frac{C(w_{n-1}w_n) + k}{C(w_{n-1}) + kV}$$

- Simple and the most popular smoothing !

STUPID BACKOFF / KATZ BACKOFF

- Longer n-grams are better but data is not always enough!
- Idea:
 - **We only back off to a lower-order n-gram if we have zero evidence for a higher-order n-gram.**

$$P_{\text{BO}}(w_n | w_{n-N+1}^{n-1}) = \begin{cases} P^*(w_n | w_{n-N+1}^{n-1}), & \text{if } C(w_{n-N+1}^n) > 0 \\ \alpha(w_{n-N+1}^{n-1}) P_{\text{BO}}(w_n | w_{n-N+2}^{n-1}), & \text{otherwise.} \end{cases}$$

INTERPOLATION SMOOTHING

- Idea:
 - **Have a combination of several n-grams models**
- E.g.:
 - The trigram probability by mixing together unigram, bigram and trigram each weighted by a λ

$$\begin{aligned}\hat{P}(w_n|w_{n-2}w_{n-1}) = & \lambda_1 P(w_n|w_{n-2}w_{n-1}) \\ & + \lambda_2 P(w_n|w_{n-1}) \\ & + \lambda_3 P(w_n)\end{aligned}$$

such that the λ s sum to 1:

$$\sum_i \lambda_i = 1$$

KNESER-NEY SMOOTHING

- One of the most commonly used and best performing n-gram smoothing methods
- Known as **absolute discounting**
- **Idea:**
 - Have the probability of the words proportional to how many different context can go before the word

$$P_{\text{AbsoluteDiscounting}}(w_i | w_{i-1}) = \frac{C(w_{i-1}w_i) - d}{\sum_v C(w_{i-1}v)} + \lambda(w_{i-1})P(w_i)$$

KNESER-NEY SMOOTHING (CONT'D)

- I can't see without my reading
- Glasses vs. Kong?
- Which one is predicted in **unigram model**?
- **Hong Kong** is more frequent word.
- A standard unigram model will assign Kong a higher probability than glasses.
- **So instead of $P(w)$, we'd like to use $P_{CONTINUATION}$**

$$P_{CONTINUATION}(w) \propto |\{v : C(vw) > 0\}|$$

- A frequent word (Kong) occurring in only one context (Hong) will have a low continuation probability

KNESER-NEY SMOOTHING (CONT'D)

- The average of d is 0.75
- λ is a normalizing constant used to distribute probability mass we've discounted

$$\lambda(w_{i-1}) = \frac{d}{\sum_v C(w_{i-1}v)} |\{w : C(w_{i-1}w) > 0\}|$$

- **Interpolated Kneser-Ney** smoothing for bigram:

$$P_{\text{KN}}(w_i|w_{i-1}) = \frac{\max(C(w_{i-1}w_i) - d, 0)}{C(w_{i-1})} + \lambda(w_{i-1})P_{\text{CONTINUATION}}(w_i)$$

Bigram count in training set	Bigram count in heldout set
0	0.0000270
1	0.448
2	1.25
3	2.24
4	3.23
5	4.21
6	5.23
7	6.21
8	7.21
9	8.26